

Regularyzacja - Lab09

Jakub Bryl

23 12 2019

Regularyzacja

Przedstawiony zbiór danych obrazuje jakość białego wina w zależności od 11 współczynników (personalne zainteresowanie na co najlepiej patrzeć przy wyborze wina w sklepie).

Wczytanie danych i normalizacja

```
data<-read_csv("winequality-white.csv")
```

```
## Parsed with column specification:
## cols(
##   `fixed acidity` = col_double(),
##   `volatile acidity` = col_double(),
##   `citric acid` = col_double(),
##   `residual sugar` = col_double(),
##   chlorides = col_double(),
##   `free sulfur dioxide` = col_double(),
##   `total sulfur dioxide` = col_double(),
##   density = col_double(),
##   pH = col_double(),
##   sulphates = col_double(),
##   alcohol = col_double(),
##   quality = col_double()
## )
```

```
summary(data)
```

```
## fixed acidity    volatile acidity    citric acid    residual sugar
## Min.   : 3.800    Min.   :0.0800    Min.   :0.0000    Min.   : 0.600
## 1st Qu.: 6.300    1st Qu.:0.2100    1st Qu.:0.2700    1st Qu.: 1.700
## Median : 6.800    Median :0.2600    Median :0.3200    Median : 5.200
## Mean   : 6.855    Mean   :0.2782    Mean   :0.3342    Mean   : 6.391
## 3rd Qu.: 7.300    3rd Qu.:0.3200    3rd Qu.:0.3900    3rd Qu.: 9.900
## Max.   :14.200    Max.   :1.1000    Max.   :1.6600    Max.   :65.800
## chlorides        free sulfur dioxide    total sulfur dioxide
## Min.   :0.00900    Min.   : 2.00         Min.   : 9.0
## 1st Qu.:0.03600    1st Qu.: 23.00         1st Qu.:108.0
## Median :0.04300    Median : 34.00         Median :134.0
## Mean   :0.04577    Mean   : 35.31         Mean   :138.4
## 3rd Qu.:0.05000    3rd Qu.: 46.00         3rd Qu.:167.0
## Max.   :0.34600    Max.   :289.00         Max.   :440.0
## density          pH          sulphates          alcohol
## Min.   :0.9871    Min.   :2.720    Min.   :0.2200    Min.   : 8.00
```

```
## 1st Qu.:0.9917 1st Qu.:3.090 1st Qu.:0.4100 1st Qu.: 9.50
## Median :0.9937 Median :3.180 Median :0.4700 Median :10.40
## Mean :0.9940 Mean :3.188 Mean :0.4898 Mean :10.51
## 3rd Qu.:0.9961 3rd Qu.:3.280 3rd Qu.:0.5500 3rd Qu.:11.40
## Max. :1.0390 Max. :3.820 Max. :1.0800 Max. :14.20
## quality
## Min. :3.000
## 1st Qu.:5.000
## Median :6.000
## Mean :5.878
## 3rd Qu.:6.000
## Max. :9.000
```

#Normalizacja

```
data_norm <- scale(data, center = TRUE, scale = TRUE) %>% as.data.frame()
summary(data_norm)
```

```
## fixed acidity    volatile acidity    citric acid    residual sugar
## Min.   :-3.61998    Min.   :-1.9668    Min.   :-2.7615    Min.   :-1.1418
## 1st Qu.: -0.65743    1st Qu.: -0.6770    1st Qu.: -0.5304    1st Qu.: -0.9250
## Median : -0.06492    Median : -0.1810    Median : -0.1173    Median : -0.2349
## Mean   : 0.00000    Mean   : 0.0000    Mean   : 0.0000    Mean   : 0.0000
## 3rd Qu.: 0.52758    3rd Qu.: 0.4143    3rd Qu.: 0.4612    3rd Qu.: 0.6917
## Max.   : 8.70422    Max.   : 8.1528    Max.   :10.9553    Max.   :11.7129
## chlorides        free sulfur dioxide    total sulfur dioxide
## Min.   :-1.6831    Min.   :-1.95848    Min.   :-3.0439
## 1st Qu.: -0.4473    1st Qu.: -0.72370    1st Qu.: -0.7144
## Median : -0.1269    Median : -0.07691    Median : -0.1026
## Mean   : 0.0000    Mean   : 0.00000    Mean   : 0.0000
## 3rd Qu.: 0.1935    3rd Qu.: 0.62867    3rd Qu.: 0.6739
## Max.   :13.7417    Max.   :14.91679    Max.   : 7.0977
## density          pH          sulphates
## Min.   :-2.31280    Min.   :-3.10109    Min.   :-2.3645
## 1st Qu.: -0.77063    1st Qu.: -0.65077    1st Qu.: -0.6996
## Median : -0.09608    Median : -0.05475    Median : -0.1739
## Mean   : 0.00000    Mean   : 0.00000    Mean   : 0.0000
## 3rd Qu.: 0.69298    3rd Qu.: 0.60750    3rd Qu.: 0.5271
## Max.   :15.02976    Max.   : 4.18365    Max.   : 5.1711
## alcohol          quality
## Min.   :-2.04309    Min.   :-3.2495
## 1st Qu.: -0.82419    1st Qu.: -0.9913
## Median : -0.09285    Median : 0.1379
## Mean   : 0.00000    Mean   : 0.0000
## 3rd Qu.: 0.71974    3rd Qu.: 0.1379
## Max.   : 2.99502    Max.   : 3.5252
```

Regresja liniowa - sprawdzenie p-val

```
model <- lm(quality ~. ,data_norm)
summary(model)
```

```
##
## Call:
## lm(formula = quality ~ ., data = data_norm)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.3299 -0.5571 -0.0428  0.5235  3.5164
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -3.668e-15  1.212e-02   0.000  1.00000
## `fixed acidity`  6.243e-02  1.989e-02   3.139  0.00171 **
## `volatile acidity` -2.120e-01  1.295e-02 -16.373 < 2e-16 ***
## `citric acid`    3.019e-03  1.309e-02   0.231  0.81759
## `residual sugar`  4.667e-01  4.311e-02  10.825 < 2e-16 ***
## chlorides       -6.100e-03  1.348e-02  -0.452  0.65097
## `free sulfur dioxide` 7.168e-02  1.621e-02   4.422 9.99e-06 ***
## `total sulfur dioxide` -1.371e-02  1.814e-02  -0.756  0.44979
## density         -5.075e-01  6.442e-02  -7.879 4.04e-15 ***
## pH              1.170e-01  1.797e-02   6.513 8.10e-11 ***
## sulphates       8.137e-02  1.294e-02   6.291 3.44e-10 ***
## alcohol         2.688e-01  3.366e-02   7.988 1.70e-15 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8484 on 4886 degrees of freedom
## Multiple R-squared:  0.2819, Adjusted R-squared:  0.2803
## F-statistic: 174.3 on 11 and 4886 DF, p-value: < 2.2e-16
```


Znalezienie optymalnego modelu

Rozdzielenie współczynników (regsubsets -> regression sub sets)

Sposób 1 - default: Forward:

```
library(leaps)


regfit.full=regsubsets(quality~.,data=data_norm, nvmax=11)
reg.summary=summary(regfit.full)
plot(reg.summary$cp,xlab="Number of Variables",ylab="Cp")
a=which.min(reg.summary$cp)
points(a,reg.summary$cp[a],pch=20,col="red")
```



Minimum wypada dla modelu zawierającego 8 wartości składowych, lecz różnica między modelem złożonym z 7 a 8 wartości jest znikoma, dlatego warto rozważyć model złożony z 7 współczynników (bez 'al sulfur dioxide').

Inna reprezentacja powyższego modelu Forward Stepwise Selection

```
plot(regfit.full,scale="Cp")
```




Selekcja modelu z użyciem zbioru walidacyjnego

Tworzymy zbiór treningowy w stosunku 70:30

```
set.seed(331)
train=sample(seq(4898),3500,replace=FALSE)
regfit.fwd=regsubsets(`quality`~.,data=data_norm[train,],nvmax=11,method="forward")
```

```
val.errors=rep(NA,11)
x.test=model.matrix(`quality`~.,data=data_norm[-train,])
for(i in 1:11){
  coefi=coef(regfit.fwd,id=i)
  pred=x.test[,names(coefi)]%*%coefi
  val.errors[i]=mean((data_norm$`quality`[-train]-pred)^2)
}
```

```
plot(sqrt(val.errors),ylab="Root MSE",ylim=c(0,3),pch=19,type="b")
points(sqrt(regfit.fwd$rss[-1]/400),col="blue",pch=19,type="b")
legend("topright",legend=c("Training","Validation"),col=c("blue","black"),pch=19)
```




Selekcja modelu z użyciem Cross-Validacji

```
predict.regsubsets=function(object,newdata,id,...){
  form=as.formula(object$call[[2]])
  mat=model.matrix(form,newdata)
  coefi=coef(object,id=id)
  mat[,names(coefi)]%*%coefi
}
```

Na poniższym wykresie został przedstawiony błąd predykcji dla modelu zawierającego X liczbę zmiennych (współczynników).


```
set.seed(11)
folds=sample(rep(1:10,length=nrow(data_norm)))
cv.errors=matrix(NA,10,11)
for(k in 1:10){
  best.fit=regsubsets(quality~.,data=data_norm[folds!=k,],nvmax=11,method="forward")
  for(i in 1:11){
    pred=predict(best.fit,data_norm[folds==k,],id=i)
    cv.errors[k,i]=mean( (data_norm$`quality`[folds==k]-pred)^2)
  }
}
rmse.cv=sqrt(apply(cv.errors,2,mean))
plot(rmse.cv,pch=19,type="b")
```




Ridge Regression

```
library(glmnet)
x=model.matrix(quality~.-1,data=data_norm)
y=data_norm$quality

#alpha=0 -> mowi ze regularyzacja l2
fit.ridge=glmnet(x,y,alpha=0)
plot(fit.ridge,xvar="lambda",label=TRUE)
```




```
cv.ridge=cv.glmnet(x,y,alpha=0)
plot(cv.ridge)
```





Lasso

```
fit.lasso=glmnet(x,y)
plot(fit.lasso,xvar="lambda",label=TRUE)
```



```
cv.lasso=cv.glmnet(x,y)
plot(cv.lasso)
```



```
coef(cv.lasso)
```

```
## 12 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept)  4.933189e-16
## `fixed acidity` -3.574513e-02
## `volatile acidity` -1.963642e-01
## `citric acid` .
## `residual sugar` 8.480340e-02
## chlorides -1.234555e-02
## `free sulfur dioxide` 4.773037e-02
## `total sulfur dioxide` .
## density .
## pH 4.333538e-03
## sulphates 2.371510e-02
## alcohol 4.608120e-01
```

Wybór parametru lambda, sprawdzamy z jaką wartością parametru lambda błąd jest najmniejszy (ponowne użycie Cross-Validacji)

```
lasso.tr=glmnet(x[train,],y[train])
lasso.tr
```

```
##
## Call:  glmnet(x = x[train, ], y = y[train])
```

```

##
##      Df      %Dev  Lambda
## 1      0 0.00000 0.44840
## 2      1 0.03303 0.40850
## 3      1 0.06044 0.37230
## 4      1 0.08321 0.33920
## 5      1 0.10210 0.30900
## 6      1 0.11780 0.28160
## 7      1 0.13080 0.25660
## 8      1 0.14160 0.23380
## 9      2 0.15670 0.21300
## 10     2 0.17270 0.19410
## 11     2 0.18610 0.17680
## 12     2 0.19710 0.16110
## 13     2 0.20630 0.14680
## 14     2 0.21400 0.13380
## 15     2 0.22030 0.12190
## 16     2 0.22550 0.11110
## 17     2 0.22990 0.10120
## 18     2 0.23350 0.09221
## 19     2 0.23650 0.08402
## 20     3 0.24090 0.07655
## 21     5 0.24710 0.06975
## 22     5 0.25250 0.06356
## 23     5 0.25700 0.05791
## 24     6 0.26080 0.05277
## 25     6 0.26430 0.04808
## 26     6 0.26730 0.04381
## 27     6 0.26970 0.03992
## 28     7 0.27180 0.03637
## 29     7 0.27350 0.03314
## 30     7 0.27500 0.03019
## 31     7 0.27620 0.02751
## 32     7 0.27720 0.02507
## 33     7 0.27800 0.02284
## 34     7 0.27870 0.02081
## 35     8 0.27920 0.01896
## 36     8 0.27970 0.01728
## 37     9 0.28080 0.01574
## 38     9 0.28290 0.01434
## 39     9 0.28460 0.01307
## 40     9 0.28600 0.01191
## 41     9 0.28720 0.01085
## 42     8 0.28800 0.00989
## 43     8 0.28850 0.00901
## 44     8 0.28890 0.00821
## 45     8 0.28920 0.00748
## 46     8 0.28950 0.00682
## 47     8 0.28970 0.00621
## 48     9 0.29000 0.00566
## 49     9 0.29050 0.00516
## 50     9 0.29090 0.00470
## 51     9 0.29130 0.00428
## 52     9 0.29150 0.00390


```

```
## 53 9 0.29180 0.00355
## 54 9 0.29200 0.00324
## 55 9 0.29220 0.00295
## 56 9 0.29230 0.00269
## 57 9 0.29240 0.00245
## 58 8 0.29250 0.00223
## 59 8 0.29260 0.00203
## 60 8 0.29270 0.00185
## 61 8 0.29270 0.00169
## 62 8 0.29280 0.00154
## 63 9 0.29280 0.00140
## 64 9 0.29290 0.00128
## 65 9 0.29290 0.00116
## 66 9 0.29290 0.00106
## 67 9 0.29290 0.00097
## 68 11 0.29300 0.00088
## 69 11 0.29300 0.00080
## 70 11 0.29300 0.00073
## 71 11 0.29300 0.00067
## 72 11 0.29300 0.00061
## 73 11 0.29300 0.00055
## 74 11 0.29300 0.00050
## 75 11 0.29300 0.00046
## 76 11 0.29300 0.00042
## 77 11 0.29300 0.00038
## 78 11 0.29300 0.00035
## 79 11 0.29300 0.00032
## 80 11 0.29300 0.00029
```

```
pred=predict(lasso.tr,x[-train,])
dim(pred)
```

```
## [1] 1398 80
```

```
rmse= sqrt(apply((y[-train]-pred)^2,2,mean))
plot(log(lasso.tr$lambda),rmse,type="b",xlab="Log(lambda)")
```



```
lam.best=lasso.tr$lambda[order(rmse)[1]]
lam.best
```

```
## [1] 0.008208603
```

```
coef(lasso.tr,s=lam.best)
```

```
## 12 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept)    0.001910617
## `fixed acidity` .
## `volatile acidity` -0.235526501
## `citric acid` .
## `residual sugar` 0.309773045
## chlorides -0.004989579
## `free sulfur dioxide` 0.050740802
## `total sulfur dioxide` .
## density -0.277472906
## pH 0.052988320
## sulphates 0.060150632
## alcohol 0.382461368
```